

# Class 7: Machine Learning 1

Erica Sanchez (PID:15787505)


Today we will start our multi-part exploration of some key machine learning methods. We will begin with clustering - finding groupings in data, and then dimensionality reduction.

## Clustering


Let's start with "k-means" cluster. The main function in base R for this is `kmeans()`.

```
# Make up some data  
hist(rnorm(100000, mean = 3))
```

**Histogram of `rnorm(1e+05, mean = 3)`**



```
tmp <- c(rnorm(30, -3), rnorm(30, +3))
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Now let's try out `kmeans()`

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	2.992899	-3.317036
2	-3.317036	2.992899

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 73.41066 73.41066
(between_SS / total_SS = 89.1 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
[1] "kmeans"
```

Q. How many points in each cluster?

```
km$size
```

```
[1] 30 30
```

Q. What component of your result object details cluster assignment/membership?

```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. What are centers/mean values of each cluster?

```
km$centers
```

```
      x      y
1  2.992899 -3.317036
2 -3.317036  2.992899
```


Q. Make a plot of your data showing your clustering results.

```
plot(x, col=km$cluster)
points(km$centers, col="purple", pch=15, cex=3)
```



Q. Run `kmeans()` again and cluster in 4 groups and plot the results.

```
km4 <- kmeans(x, centers=4)
plot(x, col=km4$cluster)
```



## Hierarchical Clustering

This form of clustering aims to reveal the structure in your data by progressively grouping points into a ever smaller number of clusters.

The main function in base R for this is called `hclust()`. This function does not take our input data directly but wants a “distance matrix” that details how (dis)similar all our input points are to each other.

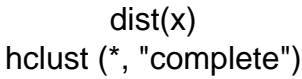
```
hc <- hclust( dist(x) )
hc
```


Call:

```
hclust(d = dist(x))
```


```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

The print out above is not very useful (unlike that from `kmeans`) but there is a useful `plot()` method.





```
plot(x, col=cutree(hc, h=6))
```



## Principal Component Analysis (PCA)

The goal of PCA is to reduce the dimensionality of a dataset down to some smaller subset of new variables (called PCs) that are useful bases for further analysis, like visualization, clustering, etc.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  4
```



## Preview the first 6 rows

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

```
head(x)
```


	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


I prefer using `x <- read.csv(url, row.names=1)` because it looks cleaner to have the column removed while it’s being read. Using the “`rownames()`” approach will remove the first unnamed column each time, so it could lead to issues when analyzing data.

Q3: Changing what optional argument in the below `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```




The so-called “pairs” plot can be useful for small datasets:

```
pairs(x, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



If a point lies on the diagonal line for a given plot, then the points have identical values.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

They do not follow the diagonal line as much as the other countries of the UK.

So the pairs plot is useful for small datasets but it can be lots of work to intercept and gets untraceable for larger datasets.

So PCA to the rescue...

The main function to do PCA in base R is called `prcomp()`. This function wants the transpose of our data in this case.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3
Standard deviation	379.8991	260.5533	1.459e-13
Proportion of Variance	0.6801	0.3199	0.000e+00
Cumulative Proportion	0.6801	1.0000	1.000e+00

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

```
pca$x
```


	PC1	PC2	PC3
Wales	-288.9534	226.36855	1.348366e-14
Scotland	-141.3603	-284.81172	4.627961e-13
N.Ireland	430.3137	58.44317	-1.280793e-13

A major PCA result viz is called a “PCA plot” (aka a score plot, biplot, PC1 vs PC2 plot, ordination plot)

Q7. Complete the code to generate a plot of PC1 vs PC2.

Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
# Plot PC1 vs PC2
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16,
      xlab="PC1", ylab="PC2")
abline(h=0, col="gray")
abline(v=0, col="gray")
```



Another important output from PCA is called “loadings” vector or the “rotation” component - this tells us how much the original variables (the foods in this case) contribute to the new PCs.

```
pca$rotation
```

	PC1	PC2	PC3
105	-0.05515951	-0.015926208	-0.114785683
245	0.05228588	-0.014247351	0.090388866
685	-0.29754711	0.017770827	-0.140679146
147	-0.08127351	0.050871661	-0.032506892
193	-0.01378426	0.095789177	-0.088241699
156	-0.03995614	0.043238659	0.060855246
720	0.38787509	0.714518360	-0.199098947
253	-0.13584223	0.144666425	-0.055102719
488	-0.24608178	0.226299860	-0.251953979
198	-0.03217825	-0.042547197	-0.091526858
360	-0.03250215	0.045390849	0.138275463
1102	-0.60208698	0.178285653	0.105684388
1472	-0.07242201	0.213840430	-0.673357649
57	-0.02896604	0.030761774	0.005505441
1374	0.21794320	-0.555250465	-0.560188267

```
375 -0.49854565 -0.110688746 -0.184979763
54 -0.03330887 -0.005704759 0.004859328
```

PCA looks to be a souper useful method for gaining some insight into high dimensional data that is difficult to examine in other ways.